

Shopify

January 18, 2022

0.1 Shopify: Summer 2022 Data Science Intern Challenge

0.1.1 By: Sunay Bhat

a. The naive Average Order Value calculation is done by taking the mean of all the order values. The issue with this dataset is **large outliers skewing the mean** - our data is not normally distributed.

b. The simplest option is to **take the median** or mode. More complicated options might include assuming a multi-modal distribution and using a cluster method to extract the AOV for each group. This notebook shows both of those methods

c. * **median** \$284.00 * **mode**: \$153.00 * **K-means** (3): * 1: \$302.58 (Majority of orders: 98.75%) * 2: \$49,213.04 (0.92% of orders) * 3: \$704,000.00 (0.34% of orders)

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
[44]: # Import and View data characteristics
df_sales = pd.read_csv('2019 Winter Data Science Intern Challenge Data Set -_
↳Sheet1.csv')
print(df_sales.head())
df_sales.nunique()
```

	order_id	shop_id	user_id	order_amount	total_items	payment_method	\
0	1	53	746	224	2	cash	
1	2	92	925	90	1	cash	
2	3	44	861	144	1	cash	
3	4	18	935	156	1	credit_card	
4	5	18	883	156	1	credit_card	

	created_at
0	2017-03-13 12:36:56
1	2017-03-03 17:38:52
2	2017-03-14 4:23:56
3	2017-03-26 12:43:37
4	2017-03-01 4:35:11

```
[44]: order_id      5000
      shop_id       100
      user_id       301
      order_amount   258
      total_items      8
      payment_method    3
      created_at     4991
      dtype: int64
```

0.1.2 a. Naive Calc issue

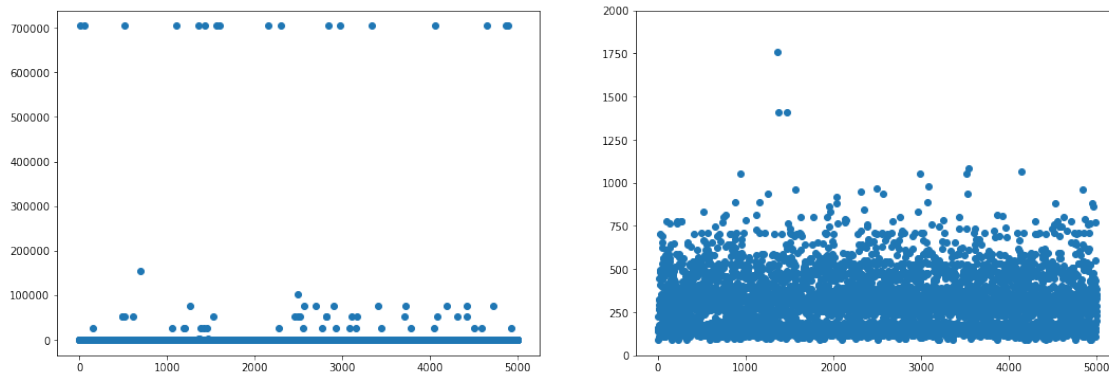
```
[45]: print('Naive AOV: ${:.2f}'.format(df_sales['order_amount'].mean()))
```

Naive AOV: \$3145.13

A quick visualization (below) shows that we have a couple modes of very large outliers that can vastly skew the mean. We would want to remove or separate our data into these modes for better analysis. A quick glance suggests 3-4 clusters, although a long-tail distribution is the most realistic model to consider.

```
[71]: # Naive Calculation is likely an average accross stores, averaged again
fig, (ax1,ax2) = plt.subplots(1,2,figsize=(18,6))
ax1.scatter(range(df_sales['order_amount'].shape[0]),df_sales['order_amount'])
ax2.scatter(range(df_sales['order_amount'].shape[0]),df_sales['order_amount'])
ax2.set_ylim([0,2000])
```

```
[71]: (0.0, 2000.0)
```



0.1.3 b and c: Alternative Calcs

```
[85]: print('Median AOV: ${:.2f}'.format(df_sales['order_amount'].median()))
      print('Mode AOV: ${:.2f}'.format(df_sales['order_amount'].mode()[0]))
```

Median AOV: \$284.00

Mode AOV: \$153.00

```
[90]: # K-means, 3 clusters
from sklearn.cluster import KMeans
km = KMeans(3, init='random', n_init=10, max_iter=300, tol=1e-04, random_state=0)
km.fit(df_sales['order_amount'].to_numpy().reshape(-1,1))
km.cluster_centers_
```

```
[90]: array([[3.02580514e+02],
           [7.04000000e+05],
           [4.92130435e+04]])
```

```
[96]: # Percent in each cluster
df_sales['kmean'] = km.labels_
df_sales['kmean'].value_counts()/50
```

```
[96]: 0    98.74
      2     0.92
      1     0.34
      Name: kmean, dtype: float64
```